

Minjune Hwang

 minjuneh@usc.edu |  <https://mj-hwang.github.io/> |  Google Scholar

EDUCATION

University of Southern California

Ph.D. in Computer Science; GPA: 4.0/4.0

Los Angeles, CA

Aug. 2024 – Present

Stanford University

M.S. in Computer Science; GPA: 3.95/4.3

Stanford, CA

Sep. 2021 – Dec. 2023

University of California, Berkeley

B.A. in Computer Science & B.A. in Statistics; GPA: 3.90/4.0

Berkeley, CA

Aug. 2017 – May. 2021

RESEARCH INTEREST

My research empowers robots to learn from humans for daily tasks by: (1) finetuning robotic foundation models with diverse feedback (e.g., demonstrations, preferences, reasoning) **to align with true user intent** and (2) improving **linguistic and spatial interpretability** of policies to allow users to **intuitively understand and guide their behavior**. In doing so I use language as a medium for human-robot interaction, developing **multimodal** and **communicative** agents that can ask and answer questions while interacting in environments.

RESEARCH EXPERIENCE

Sensing, Learning & Understanding for Robotic Manipulation (SLURM) Lab

Ph.D. Researcher, advised by Prof. Daniel Seita

Los Angeles, CA

Aug. 2024 – Present

- Building a **multimodal, communicative Vision-Language-Action** model capable of both physical (robot action), linguistic interaction (question answering/asking), and both (active perception upon question) if needed.
- Developing a structured **data augmentation** and curation method to achieve zero-shot compositional & hierarchical generalization in end-to-end VLA policies.
- Designed an **RLHF** method that uses both preference feedback and **underlying natural language reasons** to overcome causal confusion and improve policy generalization against distribution shift, with Prof. Erdem Biyik.

Stanford Vision & Learning (SVL) Lab

Graduate Research Assistant, advised by Prof. Fei-Fei Li & Prof. Jiajun Wu

Stanford, CA

Mar. 2022 – Dec. 2023

- Proposed RLHF and shared autonomy methods with **hierarchical policy learning with primitive skills** in complex manipulation tasks, first-authoring multiple papers in top conferences (CoRL, IROS).
- Designed controllers and skills for robot learning in **large-scale embodied AI simulation**, BEHAVIOR-1K.
- Researched RL algorithms with **scene graphs** for navigation & mobile manipulation with graph neural networks.

Berkeley AI Research

Research Assistant, advised by Prof. Alexandre Bayen & Prof. Laurent El Ghaoui

Berkeley, CA

Feb. 2019 – May. 2021

- Led the development of object detection and tracking model for a large-scale trajectory dataset for vehicle behavior learning and researched on safe motion planning with the dataset, under Prof. Alexandre Bayen.
- Researched extractive text summarization with sparse convex optimization, topic modeling, and RNNs.
- Developed a sparsity-invariant CNNs for adversarial attack detection via partial occlusion of images.

INDUSTRY EXPERIENCE

Amazon Robotics

Applied Scientist Intern

Westborough, MA

Sep. 2022 – Dec. 2022

- Developed ML algorithms and pipelines for object detection & tracking in warehouse human-robot collaboration.

Microsoft

Research Intern

Berkeley, CA

Jun. 2022 – Sep. 2022

- Designed an RL algorithm for offline sim-to-real transfer via reward augmentation and residual policy learning.

Apple, Special Project Group

Software Engineering Intern, Motion & Trajectory Planning

Cupertino, CA

May. 2021 – Aug. 2021

- Researched imitation learning for the warm-start of trajectory optimization of autonomous robots.
- Developed motion sampling & planning algorithms for generating safe and kinematically feasible trajectories.

SELECTED PUBLICATIONS

*: denotes equal contribution, †: equal advising

1. **Minjune Hwang**, Yigit Korkmaz, Daniel Seita[†], Erdem Biyik[†]. Causally Robust Reward Learning from Reason-Augmented Preference Feedback. *ICLR 2026*. [🔗](#) [\[poster\]](#).
2. Fangyu Wu, Dequan Wang, **Minjune Hwang**, Chenhui Hao, Jiawei Lu, Trevor Darrell, Alexandre Bayen. Decentralized Vehicle Coordination: The Berkeley DeepDrive Drone Dataset. *ICRA 2025*. [🔗](#) [\[pdf\]](#).
3. Ruohan Zhang*, Sharon Lee*, **Minjune Hwang***, Ayano Hiranaka*, Chen Wang, Wensi Ai, Jin Jie Ryan Tan, ..., Anthony Norcia, Li Fei-Fei, Jiajun Wu. NOIR: Neural Signal Operated Intelligent Robots for Everyday Activities. *CoRL 2023*. [🔗](#) [\[pdf\]](#) [🌐](#) [\[project page\]](#).
4. **Minjune Hwang***, Ayano Hiranaka*, Sharon Lee, Chen Wang, Li Fei-Fei, Jiajun Wu, Ruohan Zhang. Primitive Skill-Based Robot Learning from Human Evaluative Feedback. *IROS 2023*. [🔗](#) [\[pdf\]](#) [🌐](#) [\[project\]](#).
5. Michael Lingelbach, Chengshu Li, **Minjune Hwang**, Andrey Kurenkov, Alan Lou, Roberto Martín-Martín, Ruohan Zhang, Li Fei-Fei, Jiajun Wu. Task-Driven Graph Attention for Hierarchical Relational Object Navigation. *ICRA 2023*. [🔗](#) [\[pdf\]](#).
6. Chengshu Li*, ..., **Minjune Hwang**, ..., Silvio Savarese, Hyowon Gweon, Karen Liu, Jiajun Wu, Li Fei-Fei. BEHAVIOR-1K: A Benchmark for Embodied AI with 1,000 Everyday Activities and Realistic Simulation. *CoRL 2022*. **Best Paper Nominee** [🔗](#) [\[pdf\]](#) [🌐](#) [\[project page\]](#).
7. Michael McCoyd, Won Park, Steven Chen, Neil Shah, Ryan Roggenkemper, **Minjune Hwang**, Jason Xinyu Liu, David Wagner. Minority Reports Defense: Defending Against Adversarial Patches. *Security in Machine Learning and its Applications (SiMLA)*, 2020. **Best Paper Award** [🔗](#) [\[pdf\]](#).

TEACHING EXPERIENCE

CS 231N: Deep Learning for Computer Vision <i>Graduate Teaching Assistant</i>	Stanford, CA Mar. 2023 – Jun. 2023
EE 227BT: Convex Optimization <i>Undergraduate Teaching Assistant</i>	Berkeley, CA Aug. 2019 – Dec. 2019
CS 61A: Structure and Interpretation of Computer Programs <i>Lab Assistant</i>	Berkeley, CA Jan. 2018 – May. 2018

HONORS & FELLOWSHIPS

Viterbi School of Engineering Fellowship	Aug. 2024 – Jul. 2026
High Distinction in General Scholarship (Magna Cum Laude)	May. 2021
Summer Undergraduate Research Fellowship (SURF), UC Berkeley	May. 2020
Berkeley Undergraduate Scholarship	Aug. 2017 – May. 2021

SERVICES

- Serving/Served as a reviewer for CoRL, ICLR, ICRA, IROS, THRI, and workshops in RSS.
- Mentoring a number of undergraduate students and summer interns at USC.
- Organizing UROS, a student-run, cross-department robotics reading group and seminar series at USC.
- Serving as a PhD mentor for USC CS Undergraduate Mentoring Program.

EXPERTISES

- Robotics:** ROS, Optimal Control (LQR/LQG, MPC, MPPI), Motion Planning, Robots (Franka, Go2, XArm, WidowX)
- Robot Policy Learning:** RL (DDPG, SAC, CQL, IQL), IL (BC, ACT, Diffusion Policy), Inverse RL, PbRL, VLAs
- Computer Vision:** Object Segmentation & Tracking, Diffusion Models, Vision Transformer, Sensor Calibration, SLAM
- Libraries:** PyTorch, Tensorflow, IsaacLab, ManiSkill, Robosuite, OMPL, OpenCV, Ray, cvxopt, SageMaker, Mujoco